

Measuring the MOND Acceleration Scale at High Redshift: What Current Public Data Can and Cannot Say (the Baryon-Calibration Wall)

The authors — draft for review, 2026-09-30, version 1.1 (adds the CFG240 theorems) — AI-assisted research programme; not peer reviewed

Abstract

The programme ties the MOND acceleration scale to the vacuum, $a_0 = \kappa c \sqrt{G \rho_\Lambda}$, with $\kappa = \frac{1}{2}$ fitted, not derived. Its distinctive high-redshift prediction is a flat $a_0(z)$; the rival is $a_0 \propto H(z)$, about $\times 2.2$ at $z \approx 1.4$ and about $\times 8$ at $z \approx 5$; Λ CDM has no a_0 and enters only through an effective- a_0 proxy. We collect the programme's committed high-redshift a_0 lanes and state one result plainly: in every lane the limit that does not shrink with sample size is the absolute calibration of the baryon (gas and stellar) masses, and it is the binding one. (i) Where the discs sit ($g_{\text{bar}}/a_0 = 1.1$ to 4.4) the implied a_0 moves by -2.5 to -4.8 dex per dex of baryon mass, so a 0.05 dex error in the mass discrepancy is a factor 1.5 to 1.8 in a_0 ; in the deep regime an amplitude test measures only the product of the calibration and a_0 . (ii) No sample separates the laws: RC100 is gas-route-limited, CRISTAL is route-dependent and calibration-limited, and the $z = 2$ – 5 compilation, the ALMA [CII] rotators and the seven class-M galaxies are NOT POSSIBLE at pre-flight. (iii) The gas calibration at $z \approx 2.2$ is a prescription bracket (~ 0.2 – 0.7 dex); tracer agreement does not fix the absolute scale, and no public dynamics-independent anchor reaches 0.10 dex (the best route, an absorption-line $\alpha_{\text{[CII]}}(Z)$ relation, scatters 0.64 dex about its own fit in its own table, against the 0.2 dex it states). (iv) A within-survey lens-redshift split, which would cancel a common-mode calibration, is NOT POSSIBLE in KiDS-1000 (power 0.14 against the 9 required) and NON-DISCRIMINATING by its own MUTATE control, because a stellar-mass drift δ mimics an a_0 drift by $\delta/2$. A decisive test needs baryon errors and band ≤ 0.10 dex. With the public data on hand the question is open; nothing here is a detection either way. The decisive near-term test of the programme is a separate, $z = 0$ one: Gaia DR4 wide binaries (2 December 2026).

1 The prediction, the rival, and the size of the signal

The law and its footings. The programme's galaxy law is a MOND kernel ν with $a_0 = \kappa c \sqrt{G \rho_\Lambda}$, $\kappa = \frac{1}{2}$ fitted (not derived) and the cold component's mass still required; no dark-matter particle is added. Two footings of the local value are carried throughout: canonical $a_0 = 9.3603 \times 10^{-11}$ and alternative $1.1312 \times 10^{-10} \text{ m s}^{-2}$. The canonical footing decides unless stated; the implied ratio s^* of Section 2 is footing-relative (the alternative footing scales it by $1/1.209$, the absolute a_0 being unchanged).

Three readings of $a_0(z)$. *Flat* (the programme's distinctive law): $a_0(z) = a_0(0)$. *The rival*: $a_0(z) = a_0(0) E(z)$, with $E = H/H_0$. *The Λ CDM effective- a_0 proxy*: Λ CDM has no a_0 ; the proxy is a halo-structure scale of the record's making used as if it were $a_0(z)$ (Dutton & Macciò concentrations [DM14]), $F = 1.23, 1.77$ and 2.16 at $z = 1, 2$ and 2.5 ($+0.33$ dex at 2.5), 4.52 at $z = 4.5$ and 6.20 at $z = 5.5$. It is extrapolated above $z = 5$, and a fair Λ CDM test uses simulated galaxies (CFG222).

How large the rival's signal is, by regime. The rival's a_0 grows with z , but what a test sees depends on where the galaxies sit in g_{bar}/a_0 : in the deep regime a change of a_0 shows at half its size in the amplitude, and near the Newtonian regime it hardly shows at all.

regime	rival / flat in a_0	rival-vs-flat signal in the observable	lane
KiDS lenses, $z \approx 0.20 \rightarrow 0.40$, deep regime	small over the lever	amplitude change $+0.018$ dex (flat $+0.001$)	CFG255
NOEMA3D, $z \approx 1.4$	about $\times 2.2$	(both laws consistent there; Section 3)	CFG213
RC100, $z = 0.61$ – 2.52	—	slope of δ_{flat} on z : 0.000 if flat is true, $+0.073$ if the rival is	CFG216
$z \approx 2$ to 2.3	3.4 at $z = 2.3$; a factor-3 law difference (0.48 dex) at $z \approx 2$	implied- a_0 ratio s^*	CFG223
CRISTAL and ALPINE, $z \approx 5$	8.69 and 8.84 (the two CRISTAL points)	median δ separation 0.288 (CRISTAL) and 0.287 (ALPINE6); 0.94 dex in a_0	CFG218, CFG228, CFG224
class-M galaxies, $z \geq 2$, $g_{\text{bar}}/a_0 = 4.8$ to 157	as above	0.005 to 0.09 dex per galaxy	CFG229

The largest observable signal, 0.29 dex at $z \approx 5$ (CFG218), is smaller than the joint outer-band baryon systematic of 0.498 dex that the ALMA lane carries on the same statistic (CFG228). That comparison is the subject of this note.

2 The lever: why a small baryon error is a large a_0 error

The implied a_0 . CFG223 defines the implied ratio s^* as the multiplicative a_0 scale (relative to the footing’s local value) that sets the median over a set of galaxies of $\delta_i(s) = \log_{10}[D_i/\nu(g_{\text{bar},i}/(a_0 s))]$ to zero, with $D = g_{\text{obs}}/g_{\text{bar}}$. Flat predicts $s^* = 1$ only if the baryon calibration that set the local a_0 holds at every z . The statistical bars are small and the baryon-mass bands span the plot: RC100’s 95% intervals are a factor 2.2 to 3.8 wide, while a ± 0.15 dex shift of all baryon masses moves s^* from 0.45 to 4.04 in the lowest- z quartile and removes the solution altogether (median D below 1) in quartiles 3 and 4; ± 0.30 dex removes it at every point.

The lever. The discs sit at $g_{\text{bar}}/a_0 = 1.1$ to 4.4, where the kernel varies slowly (local slope 0.13 to 0.28 against 0.5 in the deep regime). There a 0.05 dex error in D is a factor 1.5 to 1.8 in a_0 , and $d \log_{10} s^*/d(\text{baryon dex}) = -2.5$ to -4.8 (post hoc diagnostic; the committed output prints -2.48 to -4.83 over the eight points). To put a factor-3 law difference (0.48 dex, $H(z)$ against flat at $z \approx 2$) at 2σ , the shared baryon calibration would have to be known to about 0.05 to 0.10 dex (an order-of-magnitude statement from that table).

The same lever in the other lanes. In the ALMA lane (CFG228) the lever is $\lambda = -1.8$ at the measured pooled root and -3.3 at the nominal $s^* = 1$: a 0.1 dex baryon error moves $\log_{10} s^*$ by 0.18 to 0.33 dex, against 0.66 dex between the flat and proxy expectations and 0.87 dex between flat and $H(z)$. In the class-M lane (CFG229) the one galaxy at $g_{\text{bar}}/a_0 = 4.8$ (ALESS 122.1) has $\lambda = -2.0$; for the discs at $g_{\text{bar}}/a_0 > 10$ the pre-flight lever is -56 to -70 , so a 0.05 dex baryon error would move $\log_{10} a_0$ by 2.8 to 3.5 (a factor 600 to 3,000).

The deep regime. Deep in the MOND regime $g_{\text{obs}} \simeq \sqrt{g_{\text{bar}} a_0}$, so an amplitude test there measures the product of the baryon calibration factor and a_0 : in the KiDS lensing lane (CFG255) a stellar-mass calibration drift δ between two redshift bins moves the amplitude by $\delta/2$, exactly as an a_0 drift does. Outside the deep limit the kernel’s shape does in principle separate the calibration from a_0 , so a_0 is not unidentifiable in general; but the separation rests on the curvature of ν , and the estimate becomes ill-conditioned, its error growing without bound as the sample’s largest g_{bar}/a_0 falls toward the deep regime. The formal treatment is CFG240 (below).

The conditioning, as theorems (CFG240). The lane `campaign_fresh_gravity/CFG240_calibration_wall/` formalises this in the ChainCert Lean library (50 theorems; the library goes from 281 to 331, standard axioms only). Lean certifies implications from stated premises, not any empirical fact. With one calibration factor f on g_{bar} :

- *Deep limit.* g_{obs} depends on (f, a_0) only through $f a_0$.
- *Identifiability.* For the exact P2 kernel, two distinct g_{bar} values determine (f, a_0) , and one value does not.
- *Conditioning.* The Jacobian determinant tends to zero when both points are deep.
- *A floor.* With f free, any sample obeys $\sigma(\log a_0) \geq 3\sigma/\sqrt{N}$: at a per-point error $\sigma = 0.1$ dex, a 0.1 dex a_0 needs $N > 9$; at 0.2 dex it needs $N > 36$.

Its Fisher companion (log-uniform design, P2, $N = 20$, $\sigma = 0.1$ dex, no prior on f) finds:

- A sample with $y_{\text{min}} = 0.01$ must reach $y = g_{\text{bar}}/a_0 \gtrsim 8$ (the break-even on a 0.1 dex grid of y_{max}), and one with $y_{\text{min}} = 10^{-3}$ must reach $y \gtrsim 5.5$.
- With no deep points ($y_{\text{min}} = 0.1$) it never reaches 0.1 dex: the best is 0.117.
- Even at break-even the two estimates stay correlated ($\rho \approx -0.8$): the degeneracy is broken in σ , not removed.

The high- z discs of section 2 sit at $y = 1.1$ to 4.4 with no deep points, the regime where this table gives no 0.1 dex leverage on a_0 without an external calibration of f .

What the points show (descriptive). With the caveats above (author decompositions, gas-route-limited, not a detection), flat a_0 (ratio 1) lies inside the 95% statistical interval of 7 of the 8 CFG223 points; the exception, CRISTAL at R_e on the fit route, sits 0.065 dex below its interval, so the lane’s frozen caption rule does not release the sentence “consistent with a constant a_0 ”. Inside the statistical / ± 0.15 / ± 0.30 dex band-widened intervals the counts are flat 7/8/8, proxy 5/7/7 and $H(z)$ 4/5/7. The independent-route CRISTAL points carry no information (one of them has no root in 10% of resamples).

3 What each sample gives

Each row is one committed lane; every lane froze its criteria before the numbers it scores. “Fit route” = the authors’ dynamically fitted or prior-anchored baryon masses; “independent” = SED stellar mass plus a tracer gas mass. Class S/L/M = single tracer at local prescriptions / dust or [CII] / per-galaxy optimised multi-tracer gas (the class-M band is a band on tracer agreement, not on the absolute gas scale; CFG227 correction).

lane	sample	z	N	baryon route	verdict and result in one line
CFG216, CFG217	RC100	0.61–2.52	100	fit, prior-anchored (gas scaling grows with z)	W-flat as frozen, but GAS-ROUTE-LIMITED: a differential baryon-mass systematic of about 0.15–0.25 dex over z 0.6–2.5 reverses it; it measures the differential baryon calibration, not $a_0(z)$, and its frozen significance is not to be quoted
CFG213	NOEMA3D	≈ 1.4	10	independent (measured CO)	both laws CONSISTENT in every cell, on both mass routes; no separation
CFG214	NOEMA3D outer radii	1.1–1.6	10	—	GATE-LIMITED, NON-DIAGNOSTIC AS FROZEN: 1 of 10 passes the 5% gate (beam-smeared model curves)
CFG213, CFG220	ALMA- CRISTAL	4.4–5.7	12 / 6	fit (1-dex prior on M_{bary}) / independent (dust)	fit route: rival DISFAVOURED-under, ROUTE-DEPENDENT; independent route at R_{out} : BOTH-CONSISTENT, CALIBRATION-LIMITED (the class flips at a +0.07 dex gas offset), not LOO-robust
CFG215	four samples	0.3–5.7	—	mixed	across-sample slope NON-DIAGNOSTIC by construction (route mixing; MUSE-DARK route bias +0.512 dex)
CFG227	$z = 2\text{--}5$ RAR compilation	2.0–4.4; 4.4–5.7	9 + 6	class S / L	NOT separable in any scored group; 7 of the 9 class-S sources have $g_{\text{obs}} < g_{\text{bar}}$
CFG228	six ALPINE [CII] rotators + SPT0418-47	4.2–5.5	7	class S/L (outer band ± 0.671 dex)	blind pre-flight NOT POSSIBLE; pressure term alone moves the median δ_{FLAT} from -0.467 to $+0.193$; implied a_0 UNINFORMATIVE
CFG229	class-M gold	2.024– 2.935	7	class M	pre-flight NOT POSSIBLE (Newtonian regime: signal 0.028 against SD 0.108); 6 of 7 have $D < 1$ and no root; one galaxy (ALESS 122.1) gives $s^* = 8.82$
CFG255	KiDS-1000 lenses	0.20 / 0.40	181,477	photometric M_*	NOT POSSIBLE (power 0.14 against 9) and NON-DISCRIMINATING by its MUTATE

Reading the table. *RC100* (Nestor Shachar et al. [NS23]) is the tightest sample statistically, but its baryon masses are prior-anchored with a gas scaling that grows with z . On the attack lane’s own axis the flat law’s slope is exactly zero at a differential baryon-mass change of -0.075 dex between $z \approx 0.6$ and 2.5 and the rival’s at -0.25 dex; the two sit about 0.18 dex apart, so a 3σ separation needs that differential scale known to about ± 0.06 dex on top of the statistics. Holding the gas fraction fixed with z (-0.251 on that axis) or a ULIRG-like α_{CO} (-0.211) gives the rival back; both sit at or beyond the lane’s frozen ± 0.2 dex plausible band, so they are stress tests, but the band itself is the calibration question. *CRISTAL*: on the six dust-detected discs at the outer radius the flat law’s median δ is $+0.106$ $[-0.138, +0.343]$ and the rival’s -0.203 $[-0.477, +0.029]$; the route decides it, and two discs (CRISTAL-11 and -19, 0.35 and 0.47 dex below the fit’s baryon mass) decide the route. *CFG228*: two of seven have $D < 1$, a baryon-inventory inconsistency for every law, restored by the gas band; the pooled implied $s^* = 4.33$ sits nearest the proxy’s 4.6 and that coincidence carries no information. *CFG229*: the seven sit at g_{bar}/a_0 between 4.8 (ALESS 122.1) and 157, where flat and $H(z)$ differ by 0.005 to 0.09 dex per galaxy; even with perfect baryons the pooled signal is below 2 SD, because the intrinsic scatter, velocity and inclination errors set the floor; ALESS 122.1 is one dispersion-supported galaxy whose envelope over every band corner and variant is 3.7 to 17.9 in s^* .

Forecast lanes (no law scored). CFG218: no sample is discriminating; the baryon-mass calibration needed for a 3σ -equivalent separation is 0.02 to 0.11 dex (MUSE-DARK is systematic-limited: signal 0.092 against a systematic of 0.300). CFG218 read CRISTAL as limited by statistics (about 13 discs would remove that limit); CFG219 supersedes this: “the limit is the gas calibration, not N”. CFG219: with a shared gas-calibration systematic of 0.25 dex (about 0.15 on baryons) flat against $H(z)$ never reaches 3σ at R_e for $N \leq 20$; the shared gas calibration must be known to about 0.2 dex (about 0.1–0.15 dex on the baryon mass) however many discs there are. CFG221:

the frozen pooled $z > 3.5$ rule is one-sided in practice, and a fair two-sided test needs the gas calibration to about 0.15 dex with about 50 outer-radius discs. CFG222: on RC100 the proxy sits between the flat law and $H(z)$ in every quantity (slope -0.067 against -0.029 and -0.092 ; the baryon tilt that would make each consistent is -0.18 , -0.07 and -0.25 dex), so the proxy too is calibration-bound; it is not a test of Λ CDM.

Outside this note’s lane list. The standing page records the earlier KURVS ($z \approx 1.5$) lean toward the rival as “fragile and calibration-bound”, and for MUSE-DARK the defensible wording is that “MUSE-DARK III’s a_0 rise is not recovered on SED $M_\star + H_2$; the data on disk cannot say whether the fitted masses or the SED route are biased.” Neither changes the picture here.

4 The gas calibration

What tracer agreement measures (CFG224, CFG224b). At the conversions each source states, CO, [CI] and dust gas masses agree with a pooled-scale error $K = 0.038$ at $z < 0.6$ (Stripe82, 78 galaxies) and $K = 0.128$ at $z \approx 1.2$ (14 galaxies); above $z \approx 1.6$ K is not estimated, because only three galaxies carry stated multi-tracer masses and their per-galaxy disagreements reach 0.6 dex. With the per-galaxy optimised conversions of Dunne et al. [D22] the tracer masses agree by construction, and the informative content is how the conversion factors vary with z . ACE’s dust-to-CO ratio at $z \approx 2.2$ [ACE] lies 0.67 dex below Stripe82’s relation at the frozen setting (0.21 to 0.67 depending on the metallicity slope).

What the referee found (CFG238). The independent referee reproduces 145 of 146 scored numbers, but the reading of them does not survive. (a) K is the error of a pooled mean scale, not a per-galaxy gas-mass error: the catalogue error is 0.083 to 0.111 dex per galaxy. (b) The fixed-luminosity comparison is an extrapolation: one of eight cases has a single overlapping L_{IR} bin; with sample fixed effects 3 of 9 slopes reach about 3σ ; the test has 21–52% power for a real 0.05 dex change, so “no drift larger than about 0.08 to 0.12 dex across $z = 0$ to 2.5” is what the data say at 80% power, not “no drift”. (c) ACE’s 0.2–0.7 dex is a bracket of prescriptions on an extrapolation in metallicity (span 1.09 dex across 30 cells), not a measured z -evolution. (d) A common-mode bias shared by all tracers is invisible to every diagnostic (a $+0.30$ dex shift on all tracers leaves every pair difference unchanged) and cannot be bounded from anything on disk that is independent of the dynamics under test. An earlier, more optimistic reading of CFG224 was withdrawn after this referee and is not used here.

The standing wording, quoted exactly. “multi-tracer gas is the right route, but its calibration at $z \geq 1.6$ is not shown to reach the ~ 0.1 dex a decisive $a_0(z)$ test needs; quote the gas band as a prescription bracket (~ 0.2 – 0.7 dex at $z \approx 2.2$) unless a common-mode anchor exists.”

5 No dynamics-independent anchor

A common-mode calibration needs an anchor that does not use the dynamics under test. A scoping of the public routes (paper text only, no downloads; `data_assembly/GAS_ANCHOR_SCOPING_2026-09-30.md`) reaches this bottom line: “No public route pins the common-mode gas calibration at $z \geq 1.6$ to 0.10 dex.”

The best dynamics-free route, checked. Heintz & Watson [HW20] derive $\alpha_{[\text{CI}]}(Z)$ from absorption-line column ratios in 19 GRB and QSO absorbers at z 1.962–3.376, with a stated scatter of 0.2 dex about $\log \alpha_{[\text{CI}]} = -1.13 \log(Z/Z_\odot) + 1.33$. Re-reading their Table 1: the 19 tabulated points scatter 0.64 dex rms about that relation; 7 of the 19 α rows do not follow from the paper’s own Eq. 3 applied to its columns; the tabulated α of four QSO rows is a cyclic shift of the values computed from their sources (three agree to 0.00 dex, the fourth is off by exactly 1.00 in the leading digit), and for the three GRB rows the α values follow the source CI* while the printed $N(\text{CI}^*)$ column is unexplained or matches no source value. Undoing the shift lowers the rms from 0.64 to 0.55 dex, not to the stated 0.2. Which version the authors fitted cannot be told from the paper.

The other routes. The absolute scale of Dunne et al. [D22] is assumed, not measured (a Milky-Way-like $\delta_{\text{GDR}} = 135$, i.e. $\kappa_H = 1884$, or $X_{\text{CI}} = 1.6 \times 10^{-5}$ taken from [HW20]); it constrains the relative scales of CO, [CI] and dust to 0.11–0.15 dex and says nothing on the common mode. The metallicity scale itself carries 0.14–0.18 dex of calibration systematics [C19, B23]. For one object near solar metallicity the error budget is 0.35–0.7 dex, 3 to 7 times the 0.10 dex target. Only three Dunne et al. objects with CO, [CI] and dust also have a verifiable independent metallicity. Dynamical-mass inequalities are not anchors: 68% of Stripe82 galaxies already violate $M_\star + M_{\text{CO}} + M_{\text{HI}} \leq M_{\text{dyn}}$ (an aperture mismatch), and at z 2–4 the dynamical masses are the output of the kinematic model under test (CFG238).

6 The differential route: a within-survey redshift split

A split by redshift inside one survey cancels any calibration that is common to both halves. It does not cancel a calibration that drifts with redshift, and in the deep regime a drift is indistinguishable from an a_0 drift.

KiDS-1000 by lens redshift (CFG255). On 181,477 KiDS-bright lenses (on-disk per-lens sums; nothing fetched) the extreme photo- z thirds have median z 0.20 against 0.40 (late class) and 0.23 against 0.38 (early). Over that lever the rival predicts an amplitude change of +0.018 dex and the flat law +0.001; the jackknife error on the combined amplitude is 0.038 dex; the power is $\Delta\chi^2_{\text{pred}} = 0.14$ (canonical) / 0.17 (alternative) against the 9 required, so the pre-flight said NOT POSSIBLE before the measurement. The measurement is consistent with every line: $\chi^2/14$ is 19.8 for zero and flat, 19.2 for the rival and 15.1 for the Λ CDM proxy (all $p > 0.1$); the amplitude is $+0.060 \pm 0.038$ dex, noise-level and not read as evidence either way. With the rival’s predicted ratio injected (MUTATE), flat is still not rejected ($p = 0.10$): NON-DISCRIMINATING. An independent re-run reproduced every committed results JSON byte for byte.

Why it is walled even at infinite statistics. In the deep regime a stellar-mass calibration drift δ between the bins moves the amplitude by $\delta/2$, exactly as an a_0 drift does, so the rival’s 0.018 dex needs the KiDS photometric M_* stable to better than 0.036 dex between $z \approx 0.2$ and 0.4; no source quantifies that stability. What would make the test possible: lenses reaching $z \approx 0.8$ –1.0, where the rival’s shift is 0.06–0.09 dex; a per-split amplitude error $\lesssim 0.02$ dex; and a stellar-mass drift demonstrated below about 0.05 dex. The data scoping found no public isolated-lens sample with stellar masses beyond $z = 0.5$.

The prior lane agrees. An earlier, independent two-bin KiDS lensing fit with its own interpolation function (real_research/reviews/lensing_rar/AOZ_LENSING_ZBIN_2026.md; KiDS-1000 lensing RAR of [B21]) found $a_0(\text{high})/a_0(\text{low}) = 1.309 \pm 0.210$ (statistical only) at $z_{\text{eff}} = 0.2357$ and 0.3723, against a predicted 1.083 for the rising $E(z)$ branch and 1.000 for a constant a_0 : nothing is excluded at $\geq 2\sigma$. It adds a systematic CFG255’s budget did not name: the magnitude-limited halves differ by 0.465 dex in mean $\log M_{\text{gal}}$, so a mass-dependent baryon term does not cancel between the halves. That strengthens NOT POSSIBLE; it is a null, not a detection.

7 What would decide it

The requirement. Every lane lands on the same number. The ALMA pre-flight (CFG228, post hoc arithmetic on its own mocks) shows that the six ALPINE rotators *would* separate flat from $H(z)$ if the baryon errors and the band were both ≤ 0.10 dex (signal 0.277 against $S_{\text{max}} = 0.078$ plus $2\text{SD} = 0.084$), and would not at 0.20 dex. The other lanes give the same order: 0.05 to 0.10 dex (CFG223), 0.02 to 0.11 dex (CFG218), about 0.15 dex on the gas with about 50 discs (CFG221), and a differential scale known to about ± 0.06 dex (CFG217). The standing page’s general statement: the 0.25-dex mass-scale floor caps the $z \approx 2.5$ test at 2.3σ against $a_0 \propto H(z)$ and 1.3σ against Λ CDM-native, whatever the sample size, so the calibration must reach about 0.1 dex.

What else a sample needs. From the ALMA and class-M lanes: per-galaxy optimised multi-tracer gas with stellar masses that carry a measured error; inclinations known to 5° or better (4° to 31° in the ALPINE six); rotation curves resolved beyond $2R_e$ with $\sigma(R)$ measured (the pressure term decides the sign of every δ); an intrinsic scatter below about 0.05 dex; and radii at which g_{bar}/a_0 is 1 to 3 (for the class-M discs $g_{\text{bar}}/a_0 \approx 3$ is reached only at 11 to 20 kpc, 1.3 to 8 times the stated radii). None of the public samples on disk meets these together. A common-mode gas anchor at $z \geq 1.6$ that does not use the dynamics would be needed in addition (Section 5).

The decisive near-term test is elsewhere. Gaia DR4 (2 Dec 2026) tests the programme at $z = 0$ with wide binaries; it does not measure $a_0(z)$. Candidate B (the law with hierarchical ownership) predicts $\hat{\gamma} = 1.000$ and the bare law’s floor is 1.161 (alt 1.192); with the frozen $\sigma_{\text{sys}} = 0.02$, 3σ needs about 4,300 pairs (alt 2,900) out of about 30,000 expected. It cannot separate B from Λ CDM or Newton, which both predict 1.000: for B it is a survival test, never a confirmation.

8 Limitations and what is not claimed

- This note compiles committed lanes and introduces no calculation. The implied- a_0 lever is a post hoc diagnostic of CFG223; the CFG228 requirement (≤ 0.10 dex) is post hoc arithmetic on that lane’s pre-flight mocks; CFG219’s calibration statements use a post hoc systematic-inclusive statistic, because its frozen one was blind

to a shared offset.

- All $a_0(z)$ rows rest on author decompositions (halo priors, prior-anchored baryon masses, the authors’ pressure corrections) or on our thin-disc baryon models; an implied a_0 is the scale that makes a decomposition match the kernel, not a direct measurement. Several lanes are not blind to published kinematics (CFG222, CFG223, CFG227, CFG228, CFG229 disclose this). Referees re-derived RC100 (CFG233), CRISTAL (CFG234), the compilation (CFG237) and the gas calibration (CFG238); the corrections they forced are folded in above.
- The Λ CDM proxy is not Λ CDM. Nothing here says the data favour the framework, the rival or Λ CDM; a lean is not a detection; “calibration-limited”, “NOT POSSIBLE”, “non-diagnostic” and “descriptive” carry their lanes’ meanings. $\kappa = \frac{1}{2}$ is fitted; the theory is not claimed closed; the cold mass is still required and no dark-matter particle is introduced.
- The scoping of anchors read paper text, not data; other routes may exist that it did not find, and its re-reading of [HW20] is a transcription check of a published table, not a refutation of the relation.

9 Reproducibility

Every number quoted here is re-read from a committed file by `PAPER38_audit.py`, which checks each value against its committed source at git HEAD and against the block of this source that carries its audit tag, reports how many match, and has two mutation modes (one altered expected value; one altered tex literal) that must fail. The sources are `campaign_fresh_gravity/STANDING_2026-09-29.md`, the READMEs and outputs of lanes CFG213 to CFG224, CFG227 to CFG229, CFG238 and CFG255 under `campaign_fresh_gravity/`, `data_assembly/GAS_ANCHOR_SCOPING_2026-09-30.md` and `real_research/reviews/lensing_rar/A0Z_LENSING_ZBIN_2026.md`.

References

- [ACE] The ACE survey, arXiv:2609.20926 and arXiv:2609.21040 (as cited in the CFG238 README).
- [B21] Brouwer et al. 2021, the KiDS-1000 weak-lensing radial acceleration relation (as cited in `A0Z_LENSING_ZBIN_2026.md`).
- [B23] Birkin et al. 2023, arXiv:2307.10412 (as cited in `GAS_ANCHOR_SCOPING_2026-09-30.md`).
- [C19] Coogan et al. 2019, arXiv:1901.06390 (as cited in `GAS_ANCHOR_SCOPING_2026-09-30.md`).
- [D22] Dunne et al. 2022, MNRAS 517, 962, arXiv:2208.01622 (as cited in `GAS_ANCHOR_SCOPING_2026-09-30.md`).
- [DM14] Dutton & Macciò 2014, halo concentrations (as cited in the CFG222 README).
- [HW20] Heintz & Watson 2020, ApJL 889, L7, arXiv:2001.05770 (as cited in `GAS_ANCHOR_SCOPING_2026-09-30.md`).
- [NS23] Nestor Shachar et al. 2023, the RC100 rotation-curve sample (as cited in the CFG216 README).